

Aim : Write example for following Spark RDD Actions:

**a. aggregate b. treeAggregate c. fold
d. reduce e. collect**

Solution :

1. aggregate()

```
val rdd = sc.parallelize(List(1, 2, 3, 4, 5))

// Zero value: 0, seqOp: Add elements, combOp: Combine
partitions
val result = rdd.aggregate(0)(_ + _, _ + _)

println(result) // Output: 15
```

2. treeAggregate()

```
val rdd = sc.parallelize(List(1, 2, 3, 4, 5), 3)

val result = rdd.treeAggregate(0)(_ + _, _ + _)

println(result) // Output: 15
```

3. fold()

```
val rdd = sc.parallelize(List(1, 2, 3, 4, 5))

// Zero value: 0, operation: Add elements
val result = rdd.fold(0)(_ + _)

println(result) // Output: 15
```

4. reduce()

```
val rdd = sc.parallelize(List(1, 2, 3, 4, 5))

// Reduces using a sum operation
val result = rdd.reduce(_ + _)

println(result) // Output: 15
```

5. collect()

```
val rdd = sc.parallelize(List(1, 2, 3, 4, 5))

// Collect all elements from the RDD
val result = rdd.collect()

println(result.mkString(", ")) // Output: 1, 2, 3, 4, 5
```

**Write example for following Spark RDD Actions: a. count
b. countApproxDistinct c. first d. Top
e. Min f.max**

a. count

Purpose: Counts the total number of elements in the RDD.

```
val rdd = sc.parallelize(List(10, 20, 30, 40, 50))
```

```
val totalCount = rdd.count()
```

```
println(s"Total count: $totalCount") //
```

Output: 5

b. countApproxDistinct

Purpose: Counts the approximate number of distinct elements in the RDD — more efficient for very large datasets.

```
val rdd = sc.parallelize(List(1, 2, 2, 3, 3, 3, 4, 4, 5))
```

```
val approxDistinctCount =  
rdd.countApproxDistinct()
```

```
println(s"Approximate distinct count:  
$approxDistinctCount")  
// Output: 5
```

`countApproxDistinct` uses a probabilistic algorithm — for huge data it's much faster than exact `distinct().count()`.

c. first

```
val rdd = sc.parallelize(List(100, 200, 300))
```

```
val firstElement = rdd.first()
```

```
println(s"First element: $firstElement") //
```

Output: 100

d. top

Purpose: Returns the top N elements, sorted in descending order.

```
val rdd = sc.parallelize(List(10, 70, 20, 90, 40))
```

```
val top3 = rdd.top(3) // Sorted descending
```

```
println(s"Top 3 elements: ${top3.mkString(",  
")})")  
// Output: Top 3 elements: 90, 70, 40
```

e.min

Purpose: Returns the minimum element in the RDD.

```
val rdd = sc.parallelize(List(25, 10, 30, 5,  
15))
```

```
val minValue = rdd.min()
```

```
println(s"Minimum value: $minValue")//Minimum value :5
```

Aim : Write Spark Pair RDD Functions.

A Pair RDD is an RDD where each element is a key-value pair: (K, V).

1. reduceByKey

Combine values by key with a function (e.g., sum):

```
val data = sc.parallelize(Seq(("a", 1), ("b", 2), ("a", 3)))
val result = data.reduceByKey(_ + _).collect()
// Output: Array(("a", 4), ("b", 2))
```

mapValues

Transform only the value part, keeping the key unchanged.

```
val data = sc.parallelize(Seq(("x", 2), ("y", 3)))
val squared = data.mapValues(v => v * v).collect()
// Output: Array(("x", 4), ("y", 9))
```

keys and values

Extract only keys or only values.

```
val data = sc.parallelize(Seq(("a", 1), ("b", 2)))
val keys = data.keys.collect() // Output:
Array("a", "b")
```

```
val values = data.values.collect() // Output:  
Array(1, 2)
```

sortByKey

Sort by key.

```
val data = sc.parallelize(Seq(("c", 3), ("a",  
1), ("b", 2)))  
val sorted = data.sortByKey().collect()  
// Output: Array(("a", 1), ("b", 2), ("c",  
3))
```

join

Join two Pair RDDs on key.

```
val rdd1 = sc.parallelize(Seq(("a", 1), ("b",  
2)))  
val rdd2 = sc.parallelize(Seq(("a", "x"),  
("b", "y")))  
val joined = rdd1.join(rdd2).collect()  
// Output: Array((a, (1, "x")), ("b", (2,  
"y")))
```

lookup

Return all values for a given key.

```
val data = sc.parallelize(Seq(("a", 1), ("b",  
2), ("a", 3)))  
val result = data.lookup("a")  
// Output: Seq(1, 3)
```

05/08/2025

Aim :Create two dataframes one for employee and other for dept. Perform

- a) Left anti join**
- b) Self join**
- c) Left semi join**

```
val employee = Seq(  
  (1, "Alice", 10),  
  (2, "Bob", 20),  
  (3, "Charlie", 30),  
  (4, "David", 40),  
  (5, "Eva", 50)  
).toDF("emp_id", "emp_name", "dept_id")
```

```
val dept = Seq(
```



```
(10, "HR"),  
(20, "Finance"),  
(30, "IT")  
)toDF("dept_id", "dept_name")
```

a) Left anti join

Left Anti join: Returns only the rows from the left DataFrame (e.g., **employee**) that do NOT have a match in the right DataFrame (e.g., **dept**).

"Give me employees whose department is not listed in the department table"

```
val antiJoin = employee.join(dept, Seq("dept_id"),  
"left_anti")  
antiJoin.show()
```

b) Self join

Definition:

A join of a DataFrame with itself. It is used to compare rows within the same table, often based on a common column.

"Find pairs of employees who work in the same department."

```
val e1 = employee.as("e1")  
val e2 = employee.as("e2")
```

```
val selfJoin = e1.join(e2, $"e1.dept_id" === $"e2.dept_id")  
  .select($"e1.emp_name".as("emp1"),  
    $"e2.emp_name".as("emp2"), $"e1.dept_id")  
selfJoin.show()
```

c) Left Semi Join

Returns **only the rows from the left DataFrame that have a match** in the right DataFrame — but **only columns from the left table** are returned.

"Give me employees who belong to a valid department (as per the department table)."

```
val semiJoin = employee.join(dept,  
  Seq("dept_id"), "left_semi")  
semiJoin.show()
```

- a) Create two case classes – Student and Address
- b) Create schema from these case classes

```
case class Address(city: String, state: String, pincode: String)
```

```
case class Student(id: Int, name: String, age: Int, address: Address)
```

```
val students = Seq(  
  Student(1, "Alice", 20, Address("Mumbai", "MH", "400001")),  
  Student(2, "Bob", 21, Address("Pune", "MH", "411001")),  
  Student(3, "Charlie", 22, Address("Delhi", "DL", "110001"))  
)
```

```
val studentDF = spark.createDataFrame(students)  
studentDF.printSchema()
```

3/9/2025

- a) Create a data frame with today's date and timestamp
- b) Display the hours, minutes and seconds from the timestamp

```
var df = Seq(1).toDF("id")  
df = df.withColumn("today", current_date())  
df = df.withColumn("now", current_timestamp())  
df.show()
```

b)Display the hours, minutes and seconds from the timestamp

```
var df = Seq(1).toDF("id")  
df = df.withColumn("now", current_timestamp())  
df = df.withColumn("hour", hour(col("now")))  
df = df.withColumn("minute", minute(col("now")))  
df = df.withColumn("second", second(col("now")))  
df.show(false)
```

a)For the following employee data showing name, dept and salary, perform the given operations:

**Data: ("James", "Sales", 3000),
("Michael", "Sales", 4600),
("Robert", "Sales", 4100),
("Maria", "Finance", 3000),
("James", "Sales", 3000),
("Scott", "Finance", 3300),
("Jen", "Finance", 3900),
("Jeff", "Marketing", 3000),**

**("Kumar", "Marketing", 2000),
("Saif", "Sales", 4100),
(Jason", "Sales", 9000),
("Alice", "Finance", 3700),
("Jenniffer", "Finance", 8900),
("Jenson", "Marketing", 9000)**

- a) Create a data frame for the above data**
- b) Display average salary**
- c) Display number of unique departments**
- d) Display number of employees with unique salary**

a)

```
val data = Seq(  
  ("James", "Sales", 3000),  
  ("Michael", "Sales", 4600),  
  ("Robert", "Sales", 4100),  
  ("Maria", "Finance", 3000),  
  ("James", "Sales", 3000),  
  ("Scott", "Finance", 3300),  
  ("Jen", "Finance", 3900),  
  ("Jeff", "Marketing", 3000),  
  ("Kumar", "Marketing", 2000),
```

```
("Saif", "Sales", 4100),  
("Jason", "Sales", 9000),  
("Alice", "Finance", 3700),  
("Jenniffer", "Finance", 8900),  
("Jenson", "Marketing", 9000)  
)
```

```
val df = data.toDF("name", "dept", "salary")  
df.show(false)
```

b) Display average salary

```
df.agg(avg("salary")).show()
```

c) Display number of unique departments

```
df.select("dept").distinct().count()
```

d) Display number of employees with unique salary

```
df.groupBy("salary")  
  .count()  
  .filter($"count" === 1)  
  .agg(sum("count"))  
  .show()
```

a) For the following employee data showing name, dept and salary, perform the given operations:

**Data: ("James", "Sales", 3000),
("Michael", "Sales", 4600),
("Robert", "Sales", 4100),
("Maria", "Finance", 3000),
("James", "Sales", 3000),
("Scott", "Finance", 3300),
("Jen", "Finance", 3900),
("Jeff", "Marketing", 3000),
("Kumar", "Marketing", 2000),
("Saif", "Sales", 4100),
("Jason", "Sales", 9000),
("Alice", "Finance", 3700),
("Jenniffer", "Finance", 8900),
("Jenson", "Marketing", 9000)**

- a) Create a data frame for the above data**
- b) Find the highest salary value**
- c) Find the lowest salary value**
- d) Find the standard deviation for the salary**

a)

```
val data = Seq(  
  ("James", "Sales", 3000),  
  ("Michael", "Sales", 4600),  
  ("Robert", "Sales", 4100),  
  ("Maria", "Finance", 3000),  
  ("James", "Sales", 3000),  
  ("Scott", "Finance", 3300),  
  ("Jen", "Finance", 3900),
```

```
("Jeff", "Marketing", 3000),  
("Kumar", "Marketing", 2000),  
("Saif", "Sales", 4100),  
("Jason", "Sales", 9000),  
("Alice", "Finance", 3700),  
("Jenniffer", "Finance", 8900),  
("Jenson", "Marketing", 9000)  
)
```

```
val df = data.toDF("name", "dept", "salary")  
df.show(false)
```

b) Find the highest salary

```
df.agg(max("salary")).show()
```

c) Find the lowest salary

```
df.agg(min("salary")).show()
```

d) Find the standard deviation of salary

```
df.agg(stddev("salary")).show()
```

b) Create a data frame with data that follows the below given schema

emp_id, dept, properties (a structure containing salary and location)

Return the map keys from spark SQL for this data frame


```
import org.apache.spark.sql.Row
import org.apache.spark.sql.types._
import org.apache.spark.sql.functions._
import spark.implicits._

// Schema
val schema = StructType(Seq(
  StructField("emp_id", IntegerType, false),
  StructField("dept", StringType, true),
  StructField("properties", StructType(Seq(
    StructField("salary", IntegerType, true),
    StructField("location", StringType, true)
  )))
))

//Data
val data = Seq(
  Row(1, "Sales", Row(3000, "Mumbai")),
  Row(2, "Finance", Row(4000, "Delhi")),
  Row(3, "Marketing", Row(5000, "Pune"))
)

// Create DataFrame df
val df = spark.createDataFrame(
  spark.sparkContext.parallelize(data),
  schema
)
```

```
// Convert struct -> map (df2)
val df2 = df.withColumn("properties_map", map(lit("salary"),
col("properties.salary"),lit("location"),
col("properties.location"))))
df.show(false)
```

a)

1	Create table as follows containing array and map operations		
	name	knownLanguages	properties
	James	[Java, Scala]	{hair -> black, eye -> brown}
	Michael	[Spark, Java, null]	{hair -> brown, eye -> null}
	Robert	[CSharp,]	{hair -> red, eye -> }
	Washington	null	null
	Jefferson	[]	{}

// 1. Sample data

```
val data = Seq(
  ("James",  Seq("Java", "Scala"),  Map("hair" -> "black",
"eye" -> "brown")),
  ("Michael", Seq("Spark", "Java", null), Map("hair" -> "brown",
"eye" -> null)),
  ("Robert",  Seq("CSharp", ""),    Map("hair" -> "red", "eye"
-> "")),
  ("Washington", null,              null),
  ("Jefferson", Seq.empty[String],  Map.empty[String, String])
)
```

// 2. Create DataFrame

```
val df = spark.createDataFrame(data).toDF("name",
"knownLanguages", "properties")
```

```
// 3. Show the table  
df.show(false)
```

b)

Find current timestamp and hour, Minute, second separately for today's date

```
val df = spark.range(1).select(  
  current_timestamp().alias("current_ts"),  
  hour(current_timestamp()).alias("hour"),  
  minute(current_timestamp()).alias("minute"),  
  second(current_timestamp()).alias("second")  
)  
df.show(false)
```

10/09/2025

1. a) Create a data frame containing today's date, date 2022-01-31, date 2021-03-22, date 2024-01-31, date 2023-11-11.

b) Store the date in the format MM-DD-YYYY.

c) Display the dates in the format DD/MM/YYYY

d) Find the number of months between each of the dates and today's date

Solution :

a)

```
import java.time.LocalDate
```

```
// Today's date
```

```
val today = LocalDate.now().toString // e.g. "2025-09-09"
```

```
// Create a Seq of dates
```

```
val dates = Seq(
```

```
    today,
```

```
    "2022-01-31",
```

```
    "2021-03-22",
```

```
    "2024-01-31",
```

```
    "2023-11-11"
```

```
)
```

```
// Convert to DataFrame
```

```
val df = dates.toDF("date_str")
```

```
df.show()
```

b)

```
// Convert to Date Type first
```

```
val df1 = df.withColumn("date", to_date($"date_str",  
"yyyy-MM-dd"))
```

```
// Overwrite with MM-dd-yyyy as String
```

```
val df2 = df1.withColumn("date", date_format($"date",  
"MM-dd-yyyy"))
```

```
df2.show(false)
```

c)

```
// Convert to Date Type
```

```
val df1 = df.withColumn("date", to_date($"date_str",  
"yyyy-MM-dd"))
```

```
// Display as DD/MM/YYYY
```

```
val df2 = df1.withColumn("date_DDMMYYYY",  
date_format($"date", "dd/MM/yyyy"))
```

```
df2.show(false)
```

d)

// Step 1: Convert to DateType

```
val df1 = df.withColumn("date", to_date($"date_str",  
"yyyy-MM-dd"))
```

// Step 2: Find months difference between today's date and
each dateval df2 = df1.withColumn("months_diff",
months_between(current_date(), \$"date"))

```
df2.show(false)
```

2) Refer to the employee.json file. Perform the following operations:

a) Print the names of employees above 25 years of age.

b) Print the number of employees of different ages.

employee.json(file)

[

```
{"name":"Alice", "age":25, "dept":"IT"},
```

```
{"name":"Bob", "age":30, "dept":"HR"},
```

```
{ "name": "Charlie", "age": 28, "dept": "Finance"},  
{ "name": "David", "age": 24, "dept": "IT"}  
]
```

```
val df = spark.read.option("multiline",  
"true").json("C:/Users/Lenovo/Documents/employee.json")  
df.show(false)
```

a)

```
df.filter($"age" > 25).select("name").show(false)
```

b)

```
df.groupBy("age").count().show(false)
```

3. a) Get new dates by adding 4 days, and subtracting 7 days in below dates "2020-01-02", "2023-01-15", "2025-01-30"

```
import org.apache.spark.sql.SparkSession  
import org.apache.spark.sql.functions._  
  
// Create Spark session
```

```
val spark = SparkSession.builder()  
  .appName("DateExample")  
  .master("local[*]")  
  .getOrCreate()
```

// Step 1: Create DataFrame

```
val df = spark.createDataFrame(Seq(  
  ("2020-01-02"),  
  ("2023-01-15"),  
  ("2025-01-30")  
)).toDF("date_str")
```

// Step 2: Convert to Date Type

```
val df2 = df.withColumn("date", to_date(df("date_str"),  
  "yyyy-MM-dd"))
```

// Step 3: Add 4 days

```
val df3 = df2.withColumn("date_plus_4", date_add(df2("date"),  
  4))
```


// Step 4: Subtract 7 days

```
val df_final = df3.withColumn("date_minus_7",  
date_sub(df2("date"), 7))
```

// Step 5: Show all columns

```
df_final.show(false)
```

b)

Use the Operation Read CSV file on RDD with Scala operation

// Read CSV file as RDD[String]

```
val rdd =  
spark.sparkContext.textFile("path/to/your/employee.csv")
```

// Print first 5 lines

```
rdd.take(5).foreach(println)
```